

```
import warnings
warnings.filterwarnings('ignore')

import numpy as np
import pandas as pd

housing = pd.read_csv("Housing.csv")
housing.head()
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement
0	13300000	7420	4	2	3	yes	no	no
1	12250000	8960	4	4	4	yes	no	no
2	12250000	9960	3	2	2	yes	no	no
3	12215000	7500	4	2	2	yes	no	yes
4	11410000	7420	4	1	2	yes	yes	yes

	hotwaterheating	airconditioning	parking	prefarea	furnishingstatus
0	no	yes	2	yes	furnished
1	no	yes	3	no	furnished
2	no	no	2	yes	semi-furnished
3	no	yes	3	yes	furnished
4	no	yes	2	no	furnished

```
housing.shape
```

```
(545, 13)
```

```
housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 545 entries, 0 to 544
```

```
Data columns (total 13 columns):
```

#	Column	Non-Null Count	Dtype
0	price	545 non-null	int64
1	area	545 non-null	int64
2	bedrooms	545 non-null	int64
3	bathrooms	545 non-null	int64
4	stories	545 non-null	int64
5	mainroad	545 non-null	object
6	guestroom	545 non-null	object
7	basement	545 non-null	object

```
8 hotwaterheating 545 non-null object
9 airconditioning 545 non-null object
10 parking 545 non-null int64
11 prefarea 545 non-null object
12 furnishingstatus 545 non-null object
dtypes: int64(6), object(7)
memory usage: 55.5+ KB
```

```
housing.describe()
```

	price	area	bedrooms	bathrooms	stories
count	5.450000e+02	545.000000	545.000000	545.000000	545.000000
mean	4.766729e+06	5150.541284	2.965138	1.286239	1.805505
std	1.870440e+06	2170.141023	0.738064	0.502470	0.867492
min	1.750000e+06	1650.000000	1.000000	1.000000	1.000000
25%	3.430000e+06	3600.000000	2.000000	1.000000	1.000000
50%	4.340000e+06	4600.000000	3.000000	1.000000	2.000000
75%	5.740000e+06	6360.000000	3.000000	2.000000	2.000000
max	1.330000e+07	16200.000000	6.000000	4.000000	4.000000

	parking
count	545.000000
mean	0.693578
std	0.861586
min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	3.000000

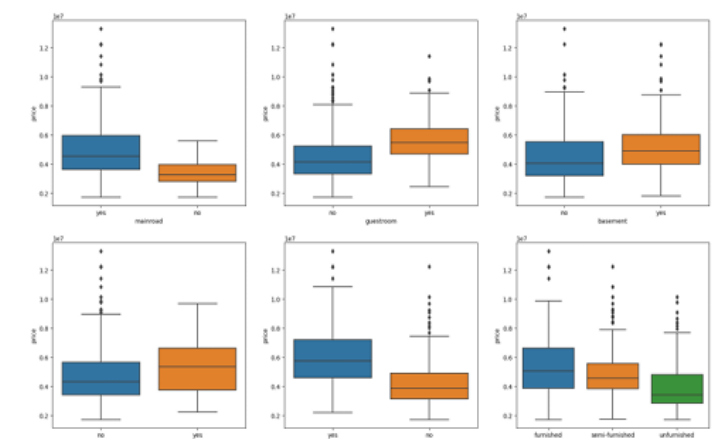
```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
sns.pairplot(housing)
plt.show()
```




```
plt.figure(figsize=(20, 12))
plt.subplot(2,3,1)
sns.boxplot(x = 'mainroad', y = 'price', data = housing)
plt.subplot(2,3,2)
sns.boxplot(x = 'guestroom', y = 'price', data = housing)
plt.subplot(2,3,3)
sns.boxplot(x = 'basement', y = 'price', data = housing)
plt.subplot(2,3,4)
sns.boxplot(x = 'hotwaterheating', y = 'price', data = housing)
plt.subplot(2,3,5)
sns.boxplot(x = 'airconditioning', y = 'price', data = housing)
plt.subplot(2,3,6)
```

```
sns.boxplot(x = 'furnishingstatus', y = 'price', data = housing)
plt.show()
```



```
plt.figure(figsize = (10, 5))
sns.boxplot(x = 'furnishingstatus', y = 'price', hue = 'airconditioning', data = housing)
plt.show()
```



```
varlist = ['mainroad', 'guestroom', 'basement', 'hotwaterheating',
'airconditioning', 'prefarea']

def binary_map(x):
    return x.map({'yes': 1, "no": 0})

housing[varlist] = housing[varlist].apply(binary_map)
housing.head()
```

	price	area	bedrooms	bathrooms	stories	mainroad
0	13300000	7420	4	2	3	1
1	12250000	8960	4	4	4	1
2	12250000	9960	3	2	2	1
3	12215000	7500	4	2	2	1
4	11410000	7420	4	1	2	1

	basement	hotwaterheating	airconditioning	parking	prefarea
0	0	0	1	2	1
1	0	0	1	3	0
2	1	0	0	2	1
3	1	0	1	3	1
4	1	0	1	2	0

```
housing.drop(['furnishingstatus'], axis = 1, inplace = True)
housing.head()
```

	price	area	bedrooms	bathrooms	stories	mainroad
0	13300000	7420	4	2	3	1
1	12250000	8960	4	4	4	1
2	12250000	9960	3	2	2	1
3	12215000	7500	4	2	2	1
4	11410000	7420	4	1	2	1

	basement	hotwaterheating	airconditioning	parking	prefarea
0	0	0	1	2	1
1	0	0	1	3	0
2	1	0	0	2	1
3	1	0	1	3	1
4	1	0	1	2	0

	semi-furnished	unfurnished
0	0	0
1	0	0
2	1	0
3	0	0
4	0	0

```
from sklearn.model_selection import train_test_split
# We specify this so that the train and test data set always have the
# same rows, respectively
np.random.seed(0)
df_train, df_test = train_test_split(housing, train_size =
0.7, test_size = 0.3, random_state = 100)

from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()

num_vars = ['area', 'bedrooms', 'bathrooms',
'stores', 'parking', 'price']
df_train[num_vars] = scaler.fit_transform(df_train[num_vars])
df_train.head()
```

	price	area	bedrooms	bathrooms	stories	mainroad
359	0.169697	0.155227	0.4	0.0	0.000000	1

```
furnishingstatus
0    furnished
1    furnished
2  semi-furnished
3    furnished
4    furnished

status = pd.get_dummies(housing['furnishingstatus'])
status.head()
```

	furnished	semi-furnished	unfurnished
0	1	0	0
1	1	0	0
2	0	1	0
3	1	0	0
4	1	0	0

```
status = pd.get_dummies(housing['furnishingstatus'], drop_first = True)
housing = pd.concat([housing, status], axis = 1)
housing.head()
```

	price	area	bedrooms	bathrooms	stories	mainroad
0	13300000	7420	4	2	3	1
1	12250000	8960	4	4	4	1
2	12250000	9960	3	2	2	1
3	12215000	7500	4	2	2	1
4	11410000	7420	4	1	2	1

	basement	hotwaterheating	airconditioning	parking	prefarea
0	0	0	1	2	1
1	0	0	1	3	0
2	1	0	0	2	1
3	1	0	1	3	1
4	1	0	1	2	0



	furnishingstatus	semi-furnished	unfurnished
0	furnished	0	0
1	furnished	0	0
2	semi-furnished	1	0
3	furnished	0	0
4	furnished	0	0

```
0
159 0.321212 0.115628 0.4 0.5 0.000000 1
1
35 0.548133 0.454417 0.4 0.5 1.000000 1
0
28 0.575758 0.538015 0.8 0.5 0.333333 1
0
```

	basement	hotwaterheating	airconditioning	parking	prefarea
359	0	0	0	0.333333	0
19	0	0	1	0.333333	1
159	1	0	1	0.000000	0
35	0	0	1	0.666667	0
28	1	1	0	0.666667	0

	semi-furnished	unfurnished
359	0	1
19	1	0
159	0	0
35	0	0
28	0	1

```
plt.figure(figsize = (16, 10))
sns.heatmap(df_train.corr(), annot = True, cmap="YlGnBu")
plt.show()
```




```
plt.figure(figsize=[6,6])
plt.scatter(df_train.area, df_train.price)
plt.show()
```

```
y_train = df_train.pop('price')
X_train = df_train

import statsmodels.api as sm
X_train_lm = sm.add_constant(X_train[['area']])
# Create a first fitted model
lr = sm.OLS(y_train, X_train_lm).fit()
# Check the parameters obtained
lr.params

const    0.126894
area     0.462192
dtype: float64

plt.scatter(X_train_lm.iloc[:, 1], y_train)
plt.plot(X_train_lm.iloc[:, 1], 0.127 + 0.462*X_train_lm.iloc[:, 1],
'r')
plt.show()
```



	coef	std err	t	P> t	[0.025
0.975]					
const	0.1269	0.013	9.853	0.000	0.102
area	0.4622	0.038	12.232	0.000	0.388
0.536					
Omnibus:		67.313	Durbin-Watson:		
2.018					
Prob(Omnibus):		0.000	Jarque-Bera (JB):		
143.063					
Skew:		0.925	Prob(JB):		
8.59e-32					
Kurtosis:		5.365	Cond. No.		
5.99					

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
X_train_lm = X_train[['area', 'bathrooms']]

import statsmodels.api as sm
X_train_lm = sm.add_constant(X_train_lm)
lr = sm.OLS(y_train, X_train_lm).fit()
lr.params

const    0.104589
area     0.398396
bathrooms 0.298374
dtype: float64

print(lr.summary())
```

```
print(lr.summary())
```

OLS Regression Results			
Dep. Variable:	price	R-squared:	0.283
Model:	OLS	Adj. R-squared:	0.281
Method:	Least Squares	F-statistic:	149.6
Date:	Thu, 17 Apr 2025	Prob (F-statistic):	3.15e-29
Time:	14:26:03	Log-Likelihood:	227.23
No. Observations:	381	AIC:	-450.5
Df Residuals:	379	BIC:	-442.6
Df Model:	1		
Covariance Type:	nonrobust		

OLS Regression Results			
Dep. Variable:	price	R-squared:	0.480
Model:	OLS	Adj. R-squared:	0.477
Method:	Least Squares	F-statistic:	174.1
Date:	Thu, 17 Apr 2025	Prob (F-statistic):	

```
bathrooms    0.259978
bedrooms     0.181863
dtype: float64
```

```
print(lr.summary())
```

OLS Regression Results					
=====					
Dep. Variable:	price	R-squared:			
0.505					
Model:	OLS	Adj. R-squared:			
0.501					
Method:	Least Squares	F-statistic:			
128.2					
Date:	Thu, 17 Apr 2025	Prob (F-statistic):			
3.12e-57					
Time:	14:31:18	Log-Likelihood:			
297.76					
No. Observations:	381	AIC:			
-587.5					
Df Residuals:	377	BIC:			
-571.7					
Df Model:	3				
Covariance Type:	nonrobust				
=====					
	coef	std err	t	P> t	[0.025
0.975]					

const	0.0414	0.018	2.292	0.022	0.006
0.077					
area	0.3922	0.032	12.279	0.000	0.329
0.455					
bathrooms	0.2600	0.026	10.033	0.000	0.209
0.311					
bedrooms	0.1819	0.041	4.396	0.000	0.101
0.263					
=====					
Omnibus:	50.037	Durbin-Watson:			
2.136					
Prob(Omnibus):	0.000	Jarque-Bera (JB):			
124.806					
Skew:	0.648	Prob(JB):			

Method:	Least Squares	F-statistic:			
60.40					
Date:	Thu, 17 Apr 2025	Prob (F-statistic):			
8.83e-83					
Time:	14:42:03	Log-Likelihood:			
381.79					
No. Observations:	381	AIC:			
-735.6					
Df Residuals:	367	BIC:			
-680.4					
Df Model:	13				
Covariance Type:	nonrobust				

=====					
	coef	std err	t	P> t	
[0.025					0.975]

const	0.0200	0.021	0.955	0.340	-
0.021	0.061				
area	0.2347	0.030	7.795	0.000	
0.175	0.294				
bedrooms	0.0467	0.037	1.267	0.206	-
0.026	0.119				
bathrooms	0.1908	0.022	8.679	0.000	
0.148	0.234				
stories	0.1085	0.019	5.661	0.000	
0.071	0.146				
mainroad	0.0504	0.014	3.520	0.000	
0.022	0.079				
guestroom	0.0304	0.014	2.233	0.026	
0.004	0.057				
basement	0.0216	0.011	1.943	0.053	-
0.000	0.043				
hotwaterheating	0.0849	0.022	3.934	0.000	
0.042	0.127				
airconditioning	0.0669	0.011	5.899	0.000	
0.045	0.089				
parking	0.0607	0.018	3.365	0.001	
0.025	0.096				
prefarea	0.0594	0.012	5.040	0.000	
0.036	0.083				
semi-furnished	0.0009	0.012	0.078	0.938	-
0.022	0.024				
unfurnished	-0.0310	0.013	-2.440	0.015	-
0.056	-0.006				

```
7.92e-28
Kurtosis:          5.487    Cond. No.
8.87
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
housing.columns
```

```
Index(['price', 'area', 'bedrooms', 'bathrooms', 'stories',
       'mainroad',
       'guestroom', 'basement', 'hotwaterheating', 'airconditioning',
       'parking', 'prefarea', 'semi-furnished', 'unfurnished'],
      dtype='object')
```

```
import statsmodels.api as sm
```

```
X_train_lm = sm.add_constant(X_train)
lr_1 = sm.OLS(y_train, X_train_lm).fit()
lr_1.params
```

```
const          0.020033
area           0.234664
bedrooms       0.046735
bathrooms      0.190823
stories        0.108516
mainroad       0.050441
guestroom      0.030428
basement       0.021595
hotwaterheating 0.084863
airconditioning 0.066881
parking        0.060735
prefarea       0.059428
semi-furnished 0.000921
unfurnished    -0.031006
dtype: float64
```

```
print(lr_1.summary())
```

OLS Regression Results					
=====					
Dep. Variable:	price	R-squared:			
0.681					
Model:	OLS	Adj. R-squared:			
0.670					
=====					
Omnibus:	93.687	Durbin-Watson:			
2.093					
Prob(Omnibus):	0.000	Jarque-Bera (JB):			
304.917					
Skew:	1.091	Prob(JB):			
6.14e-67					
Kurtosis:	6.801	Cond. No.			
14.6					
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
from statsmodels.stats.outliers_influence import
variance_inflation_factor
```

```
vif = pd.DataFrame()
vif['Features'] = X_train.columns
vif['VIF'] = [variance_inflation_factor(X_train.values, i) for i in
              range(X_train.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
```

```
vif
```

	Features	VIF
1	bedrooms	7.33
4	mainroad	6.02
0	area	4.67
3	stories	2.70
11	semi-furnished	2.19
9	parking	2.12
6	basement	2.02
12	unfurnished	1.82
8	airconditioning	1.77
2	bathrooms	1.67
10	prefarea	1.51
5	guestroom	1.47
7	hotwaterheating	1.14

Conclusion:-

Here we see that the vif score of bedrooms and mainroad is grater than 5 so, the model is not best fitted so we rejected this .